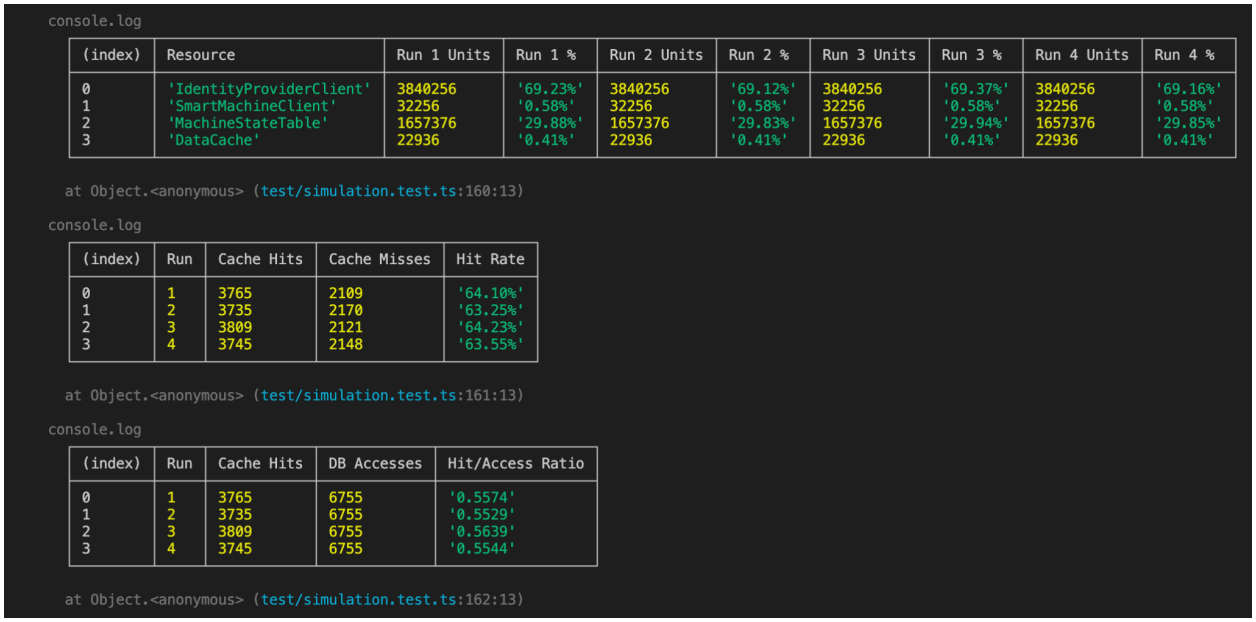

```
console.log
```

(index)	Resource	Run 1 Units	Run 1 %	Run 2 Units	Run 2 %	Run 3 Units	Run 3 %	Run 4 Units	Run 4 %
0	'IdentityProviderClient'	3840256	'69.23%'	3840256	'69.12%'	3840256	'69.37%'	3840256	'69.16%'
1	'SmartMachineClient'	32256	'0.58%'	32256	'0.58%'	32256	'0.58%'	32256	'0.58%'
2	'MachineStateTable'	1657376	'29.88%'	1657376	'29.83%'	1657376	'29.94%'	1657376	'29.85%'
3	'DataCache'	22936	'0.41%'	22936	'0.41%'	22936	'0.41%'	22936	'0.41%'

```
  at Object.<anonymous> (test/simulation.test.ts:160:13)
console.log
```

(index)	Run	Cache Hits	Cache Misses	Hit Rate
0	1	3765	2109	'64.10%'
1	2	3735	2170	'63.25%'
2	3	3809	2121	'64.23%'
3	4	3745	2148	'63.55%'

```
  at Object.<anonymous> (test/simulation.test.ts:161:13)
console.log
```

(index)	Run	Cache Hits	DB Accesses	Hit/Access Ratio
0	1	3765	6755	'0.5574'
1	2	3735	6755	'0.5529'
2	3	3809	6755	'0.5639'
3	4	3745	6755	'0.5544'

```
  at Object.<anonymous> (test/simulation.test.ts:162:13)
```

IdentityProviderClient has the highest total unit consumption (~69%), which shows that external calls are by far the most expensive part of the workflow, while SmartMachineClient and DataCache are much smaller. The cache shows about a 60% hit rate and a hit-to-DB-access ratio around 0.55, so it does reduce the database reads but still leaves room for more caching improvements.